

Start coding or [generate](#) with AI.

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Mounted at /content/drive

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
import os
```

Directory

Upload your dataset CSV (for example: nhanes_2017_2020.csv) via Colab file upload or specify path if downloaded from Google Drive

```
from google.colab import files
uploaded = files.upload() # Upload the CSV file here manually
```

↗ Choose Files demographic.csv (1).zip

- **demographic.csv (1).zip**(application/x-zip-compressed) - 376594 bytes, last modified: 7/3/2025 - 100% done
Saving demographic.csv (1).zip to demographic.csv (1).zip

download the file from kaggle

```
filename = list(uploaded.keys())[0]
print(f"Loaded file: {filename}")
```

```
df = pd.read_csv(filename)
```

↗ Loaded file: demographic.csv (1).zip

name the file

```
# Step 4: Explore data columns and preview
print("Columns in dataset:")
print(df.columns)
print("\nFirst 5 rows:")
print(df.head())
```

↗ Columns in dataset:

```
Index(['SEQN', 'SDDSRVYR', 'RIDSTATR', 'RIAGENDR', 'RIDAGEYR', 'RIDAGEMN',
       'RIDRETH1', 'RIDRETH3', 'RIDEXMON', 'RIDEXAGM', 'DMQMILIZ', 'DMQADFC',
       'DMDDBORN4', 'DMDCITZN', 'DMDYRSUS', 'DMDDEDUC3', 'DMDDEDUC2', 'DMDMARTL',
       'RIDEXPRG', 'SIALANG', 'SIAPROXY', 'SIAINTRP', 'FIALANG', 'FIAPROXY',
       'FIAINTRP', 'MIALANG', 'MIAPROXY', 'MIAINTRP', 'AIALANGA', 'DMDHHSIZ',
       'DMDFMSIZ', 'DMDHHSZA', 'DMDHHSZB', 'DMDHHSZE', 'DMDHRGND', 'DMDHRAGE',
```

```
'DMDHRBR4', 'DMDHREDU', 'DMDHRMAR', 'DMDHSEDU', 'WTINT2YR', 'WTMEC2YR',
'SDMVPSU', 'SDMVSTRA', 'INDHHIN2', 'INDFMIN2', 'INDFMPIR'],
dtype='object')
```

First 5 rows:

	SEQN	SDDSRVYR	RIDSTATR	RIAGENDR	RIDAGEYR	RIDAGEMN	RIDRETH1	\
0	73557	8	2	1	69	NaN	4	
1	73558	8	2	1	54	NaN	3	
2	73559	8	2	1	72	NaN	3	
3	73560	8	2	1	9	NaN	3	
4	73561	8	2	2	73	NaN	3	

	RIDRETH3	RIDEXMON	RIDEXAGM	...	DMDHREDU	DMDHRMAR	DMDHSEDU	\
0	4	1.0	NaN	...	3.0	4.0	NaN	
1	3	1.0	NaN	...	3.0	1.0	1.0	
2	3	2.0	NaN	...	4.0	1.0	3.0	
3	3	1.0	119.0	...	3.0	1.0	4.0	
4	3	1.0	NaN	...	5.0	1.0	5.0	

	WTINT2YR	WTMEC2YR	SDMVPSU	SDMVSTRA	INDHHIN2	INDFMIN2	INDFMPIR
0	13281.237386	13481.042095	1	112	4.0	4.0	0.84
1	23682.057386	24471.769625	1	108	7.0	7.0	1.78
2	57214.803319	57193.285376	1	109	10.0	10.0	4.51
3	55201.178592	55766.512438	2	109	9.0	9.0	2.52
4	63709.667069	65541.871229	2	116	15.0	15.0	5.00

[5 rows x 47 columns]

Select relevant columns

```
print("Columns in dataset:")
print(df.columns.tolist())
```



Columns in dataset:

```
['SEQN', 'SDDSRVYR', 'RIDSTATR', 'RIAGENDR', 'RIDAGEYR', 'RIDAGEMN', 'RIDRETH1', 'RIDRETH3', 'RIDEXMON',
```

```
...
```

```
relevant_columns = ['RIDAGEYR', 'BMXHT', 'BMXWT', 'PAQ620']
```

```
# Check which of these columns exist in the dataframe
```

```
existing_columns = [col for col in relevant_columns if col in df.columns]
```

```
print(f"Using columns: {existing_columns}")
```

```
# Select only existing columns
```

```
data = df[existing_columns].copy()
```


```
# Rename columns to standard names for convenience
```

```
rename_dict = {
    'RIDAGEYR': 'Age',
    'BMXHT': 'Height_cm',
    'BMXWT': 'Weight_kg',
    'PAQ605': 'Exercise_Level',
    'PAQ620': 'Exercise_Level' # adjust key according to your dataset
}
```

```
# Rename only the columns present
```

```
data.rename(columns={k: v for k, v in rename_dict.items() if k in data.columns}, inplace=True)
```

```
print("\nSelected data preview:")
print(data.head())
```

 Using columns: ['RIDAGEYR']

Selected data preview:

	Age
0	69
1	54
2	72
3	9
4	73


Drop rows with missing values

```
data.dropna(inplace=True)
```

Filter Exercise_Level and encode it

```
if 'Exercise_Level' in data.columns:
    # Keep only rows with Exercise_Level 1 or 2 (yes or no)
    data = data[data['Exercise_Level'].isin([1, 2])]

    # Use .loc to avoid SettingWithCopyWarning and map values
    data.loc[:, 'Exercise_Level'] = data.loc[:, 'Exercise_Level'].map({1: 2, 2: 1}) # Active=2, Inactive=
else:
    print("Warning: 'Exercise_Level' column not found in dataset. Skipping filtering and mapping.")
```

 Warning: 'Exercise_Level' column not found in dataset. Skipping filtering and mapping.

```
# Step 8: Preview cleaned data
print("\nCleaned data sample:")
print(data.head())
```



Cleaned data sample:

	Age
0	69
1	54
2	72
3	9
4	73

Define features , target and split data

```
# Step 1: Check all columns in your original dataframe
print("Columns in original dataframe:")
print(df.columns.tolist())

# Step 2: Look for weight-related columns (commonly BMXWT or similar)
weight_candidates = [col for col in df.columns if 'WT' in col.upper()]
print("Possible weight columns found:")
print(weight_candidates)
```

```

# Step 3: Define candidate columns for Age, Height, Weight, Exercise
candidate_cols = {
    'Age': 'RIDAGEYR',
    'Height_cm': 'BMXHT',
    'Weight_kg': None,          # To find below
    'Exercise_Level': None     # To find below
}

# Step 4: Select weight column from candidates found
for w_col in weight_candidates:
    if w_col in df.columns:
        candidate_cols['Weight_kg'] = w_col
        print(f"Using weight column: {w_col}")
        break

if candidate_cols['Weight_kg'] is None:
    raise KeyError("No weight column found in dataset. Cannot proceed.")

# Step 5: Find exercise column if exists
possible_exercise_cols = ['PAQ605', 'PAQ620', 'PAQ665', 'PAQ650']
for col in possible_exercise_cols:
    if col in df.columns:
        candidate_cols['Exercise_Level'] = col
        print(f"Using exercise column: {col}")
        break

# Step 6: Build list of columns that actually exist
existing_cols = [col for col in candidate_cols.values() if col is not None and col in df.columns]
print("Columns to select from dataframe:")
print(existing_cols)

# Step 7: Select columns
data = df[existing_cols].copy()

# Step 8: Rename to standard names
rename_map = {v: k for k, v in candidate_cols.items() if v is not None}
data.rename(columns=rename_map, inplace=True)

print("\nColumns after renaming:")
print(data.columns.tolist())

# Step 9: Drop rows with missing values
data.dropna(inplace=True)
print(f>Data shape after dropping NA: {data.shape}")

# Step 10: Filter Exercise_Level if it exists
if 'Exercise_Level' in data.columns:
    data = data[data['Exercise_Level'].isin([1, 2])]
    data.loc[:, 'Exercise_Level'] = data.loc[:, 'Exercise_Level'].map({1: 2, 2: 1})
else:
    print("Exercise_Level column missing; proceeding without it.")

print(f>Data shape after filtering Exercise_Level: {data.shape}")

# Step 11: Check if essential columns exist
essential_cols = ['Age', 'Weight_kg']
for col in essential_cols:
    if col not in data.columns:

```

```
raise KeyError(f"Essential column '{col}' missing from data. Cannot proceed.")
```

```
# Step 12: Define features (Height_cm optional)
features = ['Age']
if 'Height_cm' in data.columns:
    features.append('Height_cm')
if 'Exercise_Level' in data.columns:
    features.append('Exercise_Level')

print(f"Features used: {features}")

X = data[features]
y = data['Weight_kg']

print(f"Feature matrix shape: {X.shape}")
print(f"Target vector shape: {y.shape}")

# Step 13: Split data
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print(f"Training samples: {X_train.shape[0]}, Testing samples: {X_test.shape[0]}")
```

⇒ Columns in original dataframe:
 ['SEQN', 'SDDSRVYR', 'RIDSTATR', 'RIAGENDR', 'RIDAGEYR', 'RIDAGEMN', 'RIDRETH1', 'RIDRETH3', 'RIDEXMON']
 Possible weight columns found:
 ['WTINT2YR', 'WTMEC2YR']
 Using weight column: WTINT2YR
 Columns to select from dataframe:
 ['RIDAGEYR', 'WTINT2YR']

Columns after renaming:
 ['Age', 'Weight_kg']
 Data shape after dropping NA: (10175, 2)
 Exercise_Level column missing; proceeding without it.
 Data shape after filtering Exercise_Level: (10175, 2)
 Features used: ['Age']
 Feature matrix shape: (10175, 1)
 Target vector shape: (10175,)
 Training samples: 8140, Testing samples: 2035

Train Linear Regression model

```
model = LinearRegression()
model.fit(X_train, y_train)
```

⇒

▾ LinearRegression ⓘ ?
 LinearRegression()

Predict on test data

```
y_pred = model.predict(X_test)
```

Calculate Mean Squared Error

```
mse = mean_squared_error(y_test, y_pred)
print(f"\nMean Squared Error (MSE) on Test Set: {mse:.2f}")
```



Mean Squared Error (MSE) on Test Set: 717804367.81

Plot Actual vs Predicted

```
plt.figure(figsize=(7,7))
sns.scatterplot(x=y_test, y=y_pred)
plt.xlabel('Actual Weight (kg)')
plt.ylabel('Predicted Weight (kg)')
plt.title('Actual vs Predicted Weight')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.show()
```

